

Contents

AI/ML in Biomarker Discovery	1
miRNA-Based Biomarker Discovery in Alzheimer's Disease	1
6-Week Intensive Course	1
Overview	1
System Requirements	2
Part 1 — Setting Up RStudio	2
Step 1: Install R	2
Step 2: Install RStudio Desktop	2
Step 3: Configure RStudio for This Course	2
Step 4: Set Your Course Working Directory	3
Step 5: Install All R Packages (One-Time, ~15 minutes)	3
Part 2 — Setting Up VS Code for R (Alternative to RStudio)	3
Step 1: Install R (same as above)	4
Step 2: Install VS Code	4
Step 3: Install the R Extension	4
Step 4: Install the <code>languageserver</code> R Package	4
Step 5: Configure VS Code R Settings	4
Step 6: Open a Terminal and Run R	4
Part 3 — Setting Up Python (for Labs 3B, 4B, and 5B)	4
Step 1: Install Anaconda	5
Step 2: Create the Course Environment	5
Step 3: Launch JupyterLab	5
Step 4: Verify Python Packages	5
Part 4 — Project Directory Structure	5
Week-by-Week Plan	6
Week 1 — Foundations: Biology & Analytical Tools	6
Week 2 — Data Acquisition & Quality Control	6
Week 3 — Exploratory Data Analysis	6
Week 4 — Feature Selection & Classical Machine Learning	7
Week 5 — Advanced ML & External Validation	7
Week 6 — Biological Interpretation & Clinical Translation	8
Troubleshooting	8
Quick Reference: Script Execution Order	9

AI/ML in Biomarker Discovery

miRNA-Based Biomarker Discovery in Alzheimer's Disease

6-Week Intensive Course

Overview

This course teaches wet-lab biologists how to apply AI/ML techniques to discover and validate blood-based miRNA biomarkers for Alzheimer's disease (AD). No prior coding experience is required. By Week 6 you will have built, evaluated, and biologically interpreted a miRNA classifier trained on real publicly available clinical datasets.

Disease focus: Alzheimer's Disease (AD)

Biomarker class: Circulating microRNA (miRNA) from blood

Datasets: GSE120584 (serum RNA-seq, 148 samples) and GSE46579 (whole blood microarray)

Languages: R (Bioconductor) for data processing and statistics; Python (scikit-learn) for machine learning

All course materials, lecture write-ups, and R lab templates are available on GitHub:
<https://github.com/ricchasethi/ML-based-biomarker-discovery/tree/main>

System Requirements

Component	Minimum	Recommended
RAM	8 GB	16 GB
Storage	10 GB free	20 GB free
OS	Windows 10, macOS 11, Ubuntu 20.04	Any current version
Internet	Required for GEO downloads (Week 2)	Stable broadband

Part 1 — Setting Up RStudio

RStudio is the primary IDE for Weeks 1–6. All R scripts in this course are designed to be run interactively inside RStudio.

Step 1: Install R

1. Go to <https://cran.r-project.org/>
2. Click your operating system (Windows / macOS / Linux)
3. Download and run the installer for the **latest R version (4.3.0)**
 - Windows: click *base*, then *Download R x.x.x for Windows*
 - macOS: download the `.pkg` file matching your chip (Apple Silicon = `arm64`; Intel = `x86_64`)
 - Ubuntu: `sudo apt install r-base` (version from CRAN, not the default Ubuntu repo)
4. Verify installation: open a terminal and type `R --version` — you should see **R version 4.3.x** or higher

Step 2: Install RStudio Desktop

1. Go to <https://posit.co/download/rstudio-desktop/>
2. Click **Download RStudio Desktop** (free version — “RStudio Desktop Open Source License”)
3. Run the installer. On macOS, drag RStudio to your Applications folder.
4. Launch RStudio. You should see four panes:
 - **Top-left:** Script editor (where you write code)
 - **Bottom-left:** Console (where code runs and output appears)
 - **Top-right:** Environment (shows loaded variables)
 - **Bottom-right:** Files / Plots / Help

Step 3: Configure RStudio for This Course

Open RStudio and apply these settings (menu: **Tools** → **Global Options**):

Setting	Location	Recommended Value
Restore .RData on startup	General → Basic	Unchecked (prevents stale objects)
Save workspace to .RData on exit	General → Basic	Never
Default text encoding	Code → Saving	UTF-8
Soft-wrap R source files	Code → Editing	Checked

Setting	Location	Recommended Value
Rainbow parentheses	Code → Display	Checked (helps with nested code)
Show margin at column	Code → Display	80

Step 4: Set Your Course Working Directory

Every R script in this course reads and writes files relative to a single **course root folder**. Set this once at the start of every session:

```
# Replace this path with your actual course folder location
setwd("~/Documents/ML-based-biomarker-discovery")

# Verify it worked
getwd()
```

Tip: In RStudio you can also set this via Session → Set Working Directory → Choose Directory.

Step 5: Install All R Packages (One-Time, ~15 minutes)

Open miRNA/Week1_Setup_Template.R in RStudio and run **Sections 2 and 3** by pressing Ctrl+Enter (Windows/Linux) or Cmd+Enter (Mac) on each line. This installs every package needed across all 6 weeks.

To run manually:

```
# Install BiocManager first (required for all Bioconductor packages)
if (!requireNamespace("BiocManager", quietly = TRUE))
  install.packages("BiocManager")

# Bioconductor packages
BiocManager::install(c(
  "GEOquery", "DESeq2", "limma", "edgeR",
  "multiMiR", "miRBaseConverter", "clusterProfiler", "org.Hs.eg.db",
  "affy", "oligo", "sva"
), ask = FALSE, update = FALSE)

# CRAN packages
install.packages(c(
  "ggplot2", "tidyverse", "pheatmap", "reshape2", "RColorBrewer",
  "ggrepel", "gridExtra", "cluster", "factoextra", "car",
  "MASS", "pROC", "knitr", "rmarkdown"
))
```

Verify everything loaded:

```
# Paste this in the Console - if no errors appear, your environment is ready
library(GEOquery); library(DESeq2); library(limma); library(edgeR)
library(multiMiR); library(clusterProfiler); library(pheatmap)
library(ggplot2); library(cluster); library(pROC)
sessionInfo() # save this output - useful for troubleshooting
```

Part 2 — Setting Up VS Code for R (Alternative to RStudio)

If you prefer Visual Studio Code, follow these steps instead of Part 1.

Step 1: Install R (same as above)

Follow Step 1 from Part 1.

Step 2: Install VS Code

1. Go to <https://code.visualstudio.com/> and download the installer for your OS
2. Run the installer with default settings

Step 3: Install the R Extension

1. Open VS Code
2. Click the **Extensions icon** in the left sidebar (or press **Ctrl+Shift+X**)
3. Search for “**R**” and install the extension by **REditorSupport** (publisher: REditorSupport)
4. Also install the “**R Debugger**” extension by **RDebugger**

Step 4: Install the languageserver R Package

This enables autocomplete, hover documentation, and linting inside VS Code:

```
# Run this in your R terminal or RStudio Console
install.packages("languageserver")
install.packages("httpgd") # for interactive plots inside VS Code
```

Step 5: Configure VS Code R Settings

Open Settings (**Ctrl+,**), search for each setting below, and apply it:

Setting	Value
<code>r.term.windows</code>	Path to your R.exe, e.g. C:\Program Files\R\R-4.4.x\bin\R.exe
<code>r.term.mac</code>	/usr/local/bin/R (or /opt/homebrew/bin/R for Apple Silicon)
<code>r.plot.useHttpgd</code>	true (plots appear in the VS Code panel)
<code>r.lsp.enabled</code>	true
<code>r.bracketedPaste</code>	true (enables Ctrl+Enter to send selected lines to terminal)

Step 6: Open a Terminal and Run R

1. Open an integrated terminal: **Ctrl+``** (backtick)
2. Type **R** and press Enter — an interactive R session starts
3. Run code by selecting lines in the editor and pressing **Ctrl+Enter**

Note: RStudio provides a more beginner-friendly experience with a built-in file browser, environment viewer, and plot panel. If you are new to R, RStudio is strongly recommended.

Part 3 — Setting Up Python (for Labs 3B, 4B, and 5B)

Python is used for t-SNE/UMAP visualisation (Week 3) and all machine learning classifiers (Weeks 4–5). The recommended approach is Anaconda with a dedicated course environment.

Step 1: Install Anaconda

1. Go to <https://www.anaconda.com/download>
2. Download and install the **Anaconda Individual Edition** for your OS (Python 3.11+)
3. During installation on Windows: check “Add Anaconda to my PATH” (makes the next steps easier)

Step 2: Create the Course Environment

Open a terminal (macOS/Linux) or **Anaconda Prompt** (Windows):

```
conda create -n biomarker_ml python=3.11
conda activate biomarker_ml
pip install pandas numpy scipy statsmodels scikit-learn matplotlib seaborn \
    umap-learn shap xgboost lightgbm jupyterlab
```

You will need to run `conda activate biomarker_ml` at the start of every Python session.

Step 3: Launch JupyterLab

```
conda activate biomarker_ml
cd /path/to/ML-based-biomarker-discovery
jupyter lab
```

JupyterLab opens in your browser. Navigate to the `miRNA/` folder to find lab notebooks.

Step 4: Verify Python Packages

In a new Jupyter notebook, run:

```
import numpy, pandas, sklearn, matplotlib, seaborn, umap, shap, xgboost
print("All packages OK - sklearn version:", sklearn.__version__)
```

Part 4 — Project Directory Structure

Before running any script from Week 2 onward, create this folder structure inside the course root directory. The Week 2 script creates it automatically, but you can also create it manually:

```
ML-based-biomarker-discovery/
├── miRNA/                                ← all course scripts live here
│   ├── Week1_Foundations.md
│   ├── Week1_Setup_Template.R
│   ├── Week2_DataAcquisition_QC.md
│   ├── Week2_DataAcquisition_QC.R
│   ├── Week3_EDA.md
│   ├── Week3_EDA.R
│   ├── Week4_FeatureSelection_ML.md
│   ├── Week4_DE_FeatureSelection.R
│   ├── Week5_AdvancedML_Validation.md
│   ├── Week5_Validation.R
│   ├── Week6_BiologicalInterpretation.md
│   └── Week6_Interpretation.R
├── data/
│   ├── raw/                             ← GEO downloads (never modify these)
│   └── processed/                       ← clean matrices, passed between weeks
├── qc_reports/                          ← QC plots and exclusion logs
└── results/                             ← output plots and tables (Weeks 3-6)
```

Week-by-Week Plan

Week 1 — Foundations: Biology & Analytical Tools

Lecture write-up: Week1_Foundations.md

Lab script: Week1_Setup_Template.R

What you will learn: - Alzheimer's disease pathology: amyloid plaques, tau tangles, neuroinflammation, and the 15–20 year preclinical window - Why current diagnostics (CSF tau, PET imaging) are inaccessible and why blood-based biomarkers are needed - miRNA biogenesis (Drosha → Dicer → RISC), mechanism of gene silencing, and why miRNAs are stable in serum - How to distinguish microarray vs RNA-seq vs RT-qPCR measurement technologies - Key databases: GEO, miRBase v22, ADNI, AMP-AD

Lab tasks: 1. Install R 4.3.0 and RStudio (or VS Code with R extension) 2. Run Week1_Setup_Template.R Sections 2–4 to install and verify all packages 3. Complete the R basics walkthrough (vectors, data frames, ggplot2) 4. Generate your first plot: AD-relevant miRNA expression bar chart 5. Test GEOquery connectivity to NCBI servers

Deliverable: Working R environment with all packages loaded and `sessionInfo()` output saved.

Week 2 — Data Acquisition & Quality Control

Lecture write-up: Week2_DataAcquisition_QC.md

Lab script: Week2_DataAcquisition_QC.R

What you will learn: - How to navigate NCBI GEO and identify usable miRNA datasets - Downloading GSE120584 (serum RNA-seq) and GSE46579 (microarray) programmatically with GEOquery - ExpressionSet structure and metadata extraction - QC for RNA-seq: library size, detected miRNA count, RLE plots - QC for microarray: NUSE, correlation heatmap - Normalisation: DESeq2 VST (RNA-seq) and RMA (microarray) - Hemolysis detection using miR-23a/miR-451a ratio - Batch effect detection (PCA, correlation) and correction (ComBat, removeBatchEffect)

Lab tasks (run in RStudio, all sections in order): 1. Section 1: Create project directory structure 2. Sections 2–4: Download and parse both GEO datasets 3. Sections 5–7: RNA-seq QC and DESeq2 normalisation on GSE120584 4. Sections 8–9: Microarray QC and RMA normalisation on GSE46579 5. Section 10: Batch correction 6. Final: Confirm data/processed/ contains the clean .rds files

Key output files produced: - data/processed/GSE120584_expr_clean.rds — VST-normalised expression matrix - data/processed/GSE120584_metadata_clean.rds — sample metadata - data/processed/GSE120584_counts_ — raw filtered counts for DESeq2 - data/processed/GSE46579_expr_rma.rds — RMA-normalised microarray matrix

Week 3 — Exploratory Data Analysis

Lecture write-up: Week3_EDA.md

Lab script (R): Week3_EDA.R

Lab 3B (Python): Uses GSE120584_expr_vf.csv exported by Week3_EDA.R

What you will learn: - Per-miRNA statistics: coefficient of variation, IQR, zero inflation - PCA: covariance matrix intuition, scree plots, PC1/PC2 scatter, biplot, loadings table - Why t-SNE cluster sizes and distances are not interpretable; perplexity sensitivity analysis - UMAP vs t-SNE: when to use each; global structure preservation - Hierarchical clustering with Ward.D2 linkage; reading a dendrogram - Gap statistic and

silhouette width for optimal cluster number selection - Cluster purity: quantifying how well unsupervised clusters recover clinical labels - Quantifying confounder effects via partial R^2 (age and sex vs principal components) - Mahalanobis distance for outlier detection

Lab 3A tasks (RStudio, Week3_EDA.R): 1. Sections 1–4: Load data, compute stats, plot zero inflation and density 2. Section 5: Apply IQR variance filter; export CSV for Python 3. Section 6: Run PCA; generate scree plot, PC1/PC2 scatter, biplot, loadings table 4. Sections 7–10: Hierarchical clustering, gap statistic, silhouette, k-means 5. Section 11: pheatmap with Group/Sex/Age annotation tracks 6. Sections 12–13: Confounder partial R^2 and Mahalanobis outlier detection

Lab 3B tasks (JupyterLab, Python): 1. Load data/processed/GSE120584_expr_vf.csv 2. Run t-SNE at perplexity = 10, 30, 50 and compare cluster stability 3. Run UMAP at n_neighbors = 5, 15, 30; colour by Group, Age, Sex 4. Produce a 4-panel figure for the written deliverable

Key output files produced: - data/processed/GSE120584_expr_varianceFiltered.rds — input for Week 4 R script - data/processed/GSE120584_expr_vf.csv — input for Python Labs 3B and 4B - results/pca_pc1_loadings.csv, results/heatmap_top50_miRNAs.png, and others

Week 4 — Feature Selection & Classical Machine Learning

Lecture write-up: Week4_FeatureSelection_ML.md

Lab script (R): Week4_DE_FeatureSelection.R

Lab 4B (Python): Week4_ML_Classifier.ipynb — reads GSE120584_expr_forML.csv

What you will learn: - Why no single miRNA is sufficient and why a panel + ML is needed - DESeq2 workflow: negative binomial GLM, size factor normalisation, dispersion estimation, Wald test - LFC shrinkage (lfcShrink): why it matters for noisy, low-count miRNAs - Three-way disease comparisons: AD vs Control, MCI vs Control, AD vs MCI — and what each means biologically - limma-voom pipeline for the microarray validation dataset - Volcano plots and MA plots: how to read them and what artefacts to look for - Three classes of feature selection: filter (Mann-Whitney U), wrapper (RFE), embedded (LASSO, RF importance) - Building logistic regression, SVM (RBF kernel), and random forest classifiers - Model evaluation: confusion matrix, sensitivity, specificity, PPV, NPV, AUC, precision-recall curves - Why AUC alone is insufficient for imbalanced clinical datasets

Lab 4A tasks (RStudio, Week4_DE_FeatureSelection.R): 1. Sections 2–5: DESeq2 setup, run, LFC shrinkage, volcano and MA plots for all three comparisons 2. Section 6: limma pipeline on GSE46579 3. Section 7: Cross-dataset overlap analysis (DESeq2 limma) 4. Section 8: Mann-Whitney U filter feature selection 5. Section 9: Build consensus feature table 6. Section 10: Export expr_forML.csv and labels_binary.csv for Python Lab 4B

Lab 4B tasks (JupyterLab, Week4_ML_Classifier.ipynb): 1. Load exported feature matrix and binary labels 2. Apply Mann-Whitney filter; perform 80/20 stratified train/test split 3. Fit logistic regression, SVM (with grid search), and random forest 4. Generate confusion matrices, ROC curves, precision-recall curves 5. Run 5-fold cross-validation; report mean AUC $\pm$ 95% CI per model 6. Extract random forest feature importances; look up top 3 in miRTarBase

Key output files produced: - results/de_results_deseq2_AD_vs_Control.csv (and MCI vs Control, AD vs MCI) - results/consensus_features_Week4.csv — starting point for Week 5 nested CV - data/processed/GSE120584_expr_forML.csv, data/processed/GSE120584_labels_binary.csv

Week 5 — Advanced ML & External Validation

Lecture write-up: Week5_AdvancedML_Validation.md

Lab script (R): Week5_Validation.R

Lab 5B (Python): Reads harmonized matrices exported by Week5_Validation.R

What you will learn: - Bias-variance tradeoff and the five most common forms of data leakage in bioinformatics - Nested cross-validation: why a single train/test split gives biased AUC when hyperparameters are tuned - XGBoost and LightGBM: gradient boosting intuition; key hyperparameters (learning rate, n_estimators, early stopping) - SHAP (Shapley Additive Explanations): interpreting individual-sample predictions from any model - Three-class classification (AD / MCI / Control): multiclass metrics and confusion matrices - Cross-platform harmonization: miRBaseConverter, MIMAT accession mapping, z-score standardisation - DeLong's test for AUC comparison between models (pROC package) - Calibration plots and Brier score: does the model's probability output mean what it says? - TRIPOD and STARD reporting guidelines for biomarker ML studies

Lab 5A tasks (RStudio, Week5_Validation.R): 1. Section 2: Load both datasets; harmonize miRNA names via miRBaseConverter 2. Section 3: Find feature intersection between platforms 3. Section 4: Apply per-dataset z-score standardisation 4. Section 5: Export harmonized matrices for Python Lab 5B 5. Section 6: DeLong AUC comparison test 6. Section 7: Calibration plot (manual binning and Brier score) 7. Section 8: Compile model results summary table

Lab 5B tasks (JupyterLab, Python): 1. Load harmonized matrices from Week5_Validation.R 2. Implement nested cross-validation for logistic regression and random forest 3. Train XGBoost; tune with early stopping 4. Compute SHAP values; generate summary plot and beeswarm plot 5. Apply the GSE120584-trained model to GSE46579 as external validation 6. Report AUC with 95% bootstrapped confidence intervals

Key output files produced: - data/processed/GSE120584_harmonized.rds, data/processed/GSE46579_harmonized.rds
- results/calibration_plot.png, results/model_results_summary.csv

Week 6 — Biological Interpretation & Clinical Translation

Lecture write-up: Week6_BiologicalInterpretation.md

Lab script: Week6_Interpretation.R

What you will learn: - How to integrate SHAP importance scores with DE fold changes into a composite biomarker ranking - Querying 14 miRNA–target databases simultaneously with multiMiR; filtering for validated interactions - KEGG pathway enrichment analysis (does the AD pathway hsa05010 appear enriched?) - GO Biological Process enrichment and redundancy removal with clusterProfiler - Generating the “paper figure”: a forest plot combining log2FC, SHAP importance, and key targets - Regulatory considerations for a blood-based miRNA diagnostic: FDA IVD pathway overview

Lab tasks (RStudio, Week6_Interpretation.R, all sections in order): 1. Section 2: Load SHAP results and DE tables; build composite ranking score; select top 15 miRNAs 2. Section 3: multiMiR target query; filter for validated targets; cross-reference known AD genes 3. Section 4: KEGG enrichment — confirm or deny hsa05010 enrichment 4. Section 5: GO enrichment with term simplification 5. Section 6: Generate forest plot (the “main figure”) 6. Section 7: Print the full 6-week pipeline summary table

Key output files produced (in results/Week6/): - composite_biomarker_ranking.csv, validated_targets.csv
- kegg_enrichment.csv, go_enrichment_simplified.csv - biomarker_forest_plot.png

Troubleshooting

“there is no package called ‘X’”

Run `BiocManager::install("X")` (Bioconductor) or `install.packages("X")` (CRAN), then try again.

“package ‘X’ was built under R version Y.Z”

This warning is usually harmless. Continue unless you see an actual error downstream.

“Error in getGEO ... could not resolve host”

No internet connection, or NCBI servers are temporarily down. Check your connection and try again in a few minutes. On institutional networks a VPN may be required.

“‘BiocManager’ is not available for R version X.Y.Z”

Your R version is below 4.3.0. Download the latest R from <https://cran.r-project.org/> and reinstall RStudio after upgrading.

DESeq2 or limma gives different results than expected

Confirm your working directory is the course root and that all **data/processed/** files from the previous week exist. Each script depends on the outputs of all prior scripts.

Jupyter notebook cannot import umap or shap

Confirm your **biomarker_ml** conda environment is activated before launching JupyterLab: `conda activate biomarker_ml && jupyter lab`.

Quick Reference: Script Execution Order

Week 1: Week1_Setup_Template.R	← one-time package installation
Week 2: Week2_DataAcquisition_QC.R	← downloads data; run with internet access
Week 3: Week3_EDA.R	← then open Python for Lab 3B
Week 4: Week4_DE_FeatureSelection.R	← then open Week4_ML_Classifier.ipynb
Week 5: Week5_Validation.R	← then open Python for Lab 5B
Week 6: Week6_Interpretation.R	← final week; reads all prior outputs

Each script ends with a summary printout listing every file it produced and what to open next.